

471 Midterm

Zhoumengdi Wang zxw534

1.Preprocessing data/Perform exploratory analyses

(1) What is missing data?

The first thing we need to do is preprocessing the data. Now we need to take care of the missing data. `Hmisc::describe(census_train)`, this function can help us find the missing data. I use `Hmisc::describe` to view the whole data.

According to the form, there is no missing data, but I found:

```
-----
workclass      occupation      native.country
  n missing distinct      n missing distinct      n missing distinct
25000         0         9      25000         0         15      25000         0         42
? (1404, 0.056), Federal-gov (?) ? (1411, 0.056), Adm-clerical  lowest : ?
                                   highest: Thailand
```

We can see that there are three variables has the value "?". Although "?" isn't regarded as missing data, but "?" means nothing in real. In fact, "?" is the missing data.

(2) How to deal with the missing data

With the missing data, there are some ways to deal with:

1. delete
2. use the highest frequency value to replace the missing data
3. use the relationships between variables

I will combine 2 and 3 to deal with the missing data. The first class which has the missing data is workclass. I think the workclass is influenced by occupation. Because different workclass is based on what kind of job you have. So my first step is dealing with the missing data in occupation. For the missing data in occupation, I would like to use highest frequency value to replace the missing data. For occupation, the highest frequency value is "Prof-specialty" so I use "Prof-specialty" to replace the missing data in occupation.

After dealing with the missing data in occupation, I need to find the relationship between the occupation and workclass. Show the table of these two variables:

| | Federal-gov | Local-government | Never-worked | Private | Self-employed-incorporated | Self-employed-not-incorporated | State-government | Without-pay |
|--------------------|-------------|------------------|--------------|---------|----------------------------|--------------------------------|------------------|-------------|
| Adm-clerical | 240 | 220 | 0 | 2146 | 19 | 37 | 186 | 2 |
| Armed-Forces | 6 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| Craft-repair | 47 | 112 | 0 | 2418 | 82 | 412 | 38 | 1 |
| Exec-managerial | 133 | 162 | 0 | 2064 | 298 | 290 | 141 | 0 |
| Farming-fishing | 5 | 21 | 0 | 357 | 33 | 333 | 13 | 4 |
| Handlers-cleaners | 17 | 34 | 0 | 967 | 1 | 11 | 6 | 0 |
| Machine-operatives | 11 | 9 | 0 | 1436 | 8 | 28 | 11 | 1 |
| Other-service | 27 | 142 | 0 | 2051 | 19 | 141 | 97 | 1 |
| Priv-household | 0 | 0 | 0 | 108 | 0 | 0 | 0 | 0 |
| Prof-specialty | 134 | 532 | 0 | 1740 | 120 | 269 | 316 | 0 |
| Protective | 22 | 228 | 0 | 142 | 2 | 5 | 84 | 0 |
| Sales | 11 | 5 | 0 | 2202 | 219 | 284 | 9 | 0 |
| Tech-support | 52 | 26 | 0 | 542 | 3 | 18 | 44 | 0 |
| Transport | 18 | 91 | 0 | 963 | 23 | 88 | 32 | 1 |

We can find that when the occupation is “Prof-specialty”, the highest frequency value of workclass is “Private”. I use private to replace the missing data in workclass.

There is one variable which has missing date, the native_country. Because it is a US census, the native country of most people in this census is US, so the highest frequency value of native country is United_States.

After dealing with the missing data and checking there is no left missing data, group two part data.

```
census_train_new = rbind(census_train_completed, census_train_missing)
```

(3) Check the variables

Close definition variables

There are two pairs variables, which has close definition, education and education.num, marital.status and relationship. Here are the tables of these two pairs:

education and education.num,

| | 10th | 11th | 12th | 1st-4th | 5th-6th | 7th-8th | 9th | Assoc-ac | Assoc-v | Bachelor | Doctorat | HS-grad | Masters | Preschod | Prof-sch | Some-col |
|----|------|------|------|---------|---------|---------|-----|----------|---------|----------|----------|---------|---------|----------|----------|----------|
| 1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 35 | 0 | 0 |
| 2 | 0 | 0 | 0 | 129 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 3 | 0 | 0 | 0 | 0 | 255 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 4 | 0 | 0 | 0 | 0 | 0 | 498 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 5 | 0 | 0 | 0 | 0 | 0 | 0 | 402 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 6 | 726 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 7 | 0 | 897 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 8 | 0 | 0 | 355 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 9 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 8064 | 0 | 0 | 0 | 0 |
| 10 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 5585 |
| 11 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 1062 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 12 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 822 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 13 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 4105 | 0 | 0 | 0 | 0 | 0 | 0 |
| 14 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 1304 | 0 | 0 | 0 | 0 |
| 15 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 450 | 0 |
| 16 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 311 | 0 | 0 | 0 | 0 | 0 |







The table shows that each class of education.num matches a education level, so I think the education should be dropped.

marital.status and relationship

| | Husband | Not-in-f | Other-re | Own-chil | Unmarrie | Wife |
|----------|---------|----------|----------|----------|----------|------|
| Divorced | 0 | 1843 | 90 | 256 | 1236 | 0 |
| Married- | 7 | 0 | 1 | 1 | 0 | 9 |
| Married- | 10110 | 16 | 88 | 64 | 0 | 1215 |
| Married- | 0 | 163 | 26 | 31 | 107 | 0 |
| Never-ma | 0 | 3607 | 487 | 3446 | 656 | 0 |
| Separate | 0 | 321 | 44 | 78 | 347 | 0 |
| Widowed | 0 | 420 | 40 | 11 | 280 | 0 |






I think marital.status is more clear than the relationship, so I decide to drop relationship. And I will combine “Married-AF-spouse” and “Married-civ-spouse” which seem like similar into a single category called “Married”.

Check quantitative data




Distribution of most quantitative data is fine, however, capital.gain and capital.loss of most people is zero. I may delete them in adjusting models.

Watch qualitative data



native



For native.country, almost all people's native.country is United-States. So I will combine other countries into a new label "other country"

```
## [1] "Other Country" "United-States"
```

Check collinearity to find if there is some relationships between quantitative variables.

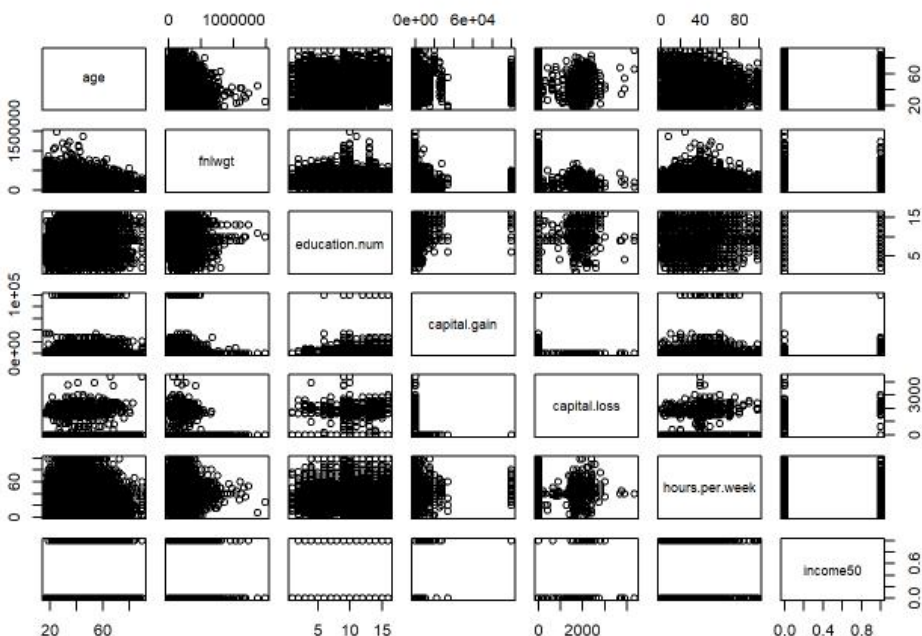
```
cor(census_train_new_quantitative)
```

```
##           age           fnlwgt education.num capital.gain
## age           1.00000000 -0.0754229958    0.03745462  0.0803364007
## fnlwgt        -0.07542300  1.0000000000   -0.04596356 -0.0004977309
## education.num  0.03745462 -0.0459635576    1.00000000  0.1299508553
## capital.gain   0.08033640 -0.0004977309    0.12995086  1.0000000000
## capital.loss   0.06110263 -0.0109322145    0.08314796 -0.0318820464
## hours.per.week 0.06864841 -0.0153756410    0.15357666  0.0850860033
##
##           capital.loss hours.per.week
## age           0.06110263    0.06864841
## fnlwgt        -0.01093221   -0.01537564
## education.num  0.08314796    0.15357666
## capital.gain   -0.03188205    0.08508600
## capital.loss   1.00000000    0.05688347
## hours.per.week 0.05688347    1.00000000
```

Most variables do not have strong collinearity, but the collinearity values of education.num-capital.gain and education.num-hours.per.week is higher than 0.1. According to ISLR, I can delete some variables which have collinearity. But, I don't want to delete education.num-hours.per.week. Because I just choose education.num rather than education, and hours.per.week is obviously related to income. How much time you work, how much money you can get.

check variables and income

Find whether there is some variable can determinate income directly or has some determinative role.



No quantitative variable is the determinative role of income.

```
table(census_train_new$workclass, census_train_new$income)
```

```
##
##
##      Federal-gov      <=50K >50K
##      Local-gov       1141   465
##      Never-worked      7     0
##      Private        14888  3963
##      Self-emp-inc      381   475
##      Self-emp-not-inc  1396  547
##      State-gov        732   261
##      Without-pay      10     0
```

```
table(census_train_new$marital.status, census_train_new$income)
```

```
##
##
##      Divorced      <=50K >50K
##      Married      6361  5150
##      Married-spouse-absent  302   25
##      Never-married  7844  352
##      Separated      735   55
##      Widowed       690   61
```

```
table(census_train_new$occupation, census_train_new$income)
```

```
##
##
##      Adm-clerical      <=50K >50K
##      Armed-Forces      5       1
##      Craft-repair     2441   719
##      Exec-managerial   1628  1518
##      Farming-fishing    676    94
##      Handlers-cleaners  981    70
##      Machine-op-inspct 1354   180
##      Other-service     2435   101
##      Priv-house-serv    113     1
##      Prof-specialty    1771  1415
##      Protective-serv    330   157
##      Sales             2043   743
##      Tech-support      486   211
##      Transport-moving   977   254
##      Prof-specialty    1259   152
```

```
table(census_train_new$race, census_train_new$income)
```

```
##
##
##      Amer-Indian-Eskimo  214    25
##      Asian-Pac-Islander  592   214
##      Black              2077   292
##      Other               197    19
##      White              15922  5448
```

```
table(census_train_new$sex, census_train_new$income)
```

```
##
##
##      Female      <=50K >50K
##      Male       11692  5095
```

```
table(census_train_new$native.country, census_train_new$income)
```

```
##
##
##      Other Country      <=50K >50K
##      United-States    1773    396
##      United-States    17229  5602
```

No qualitative variable has the determinative role of income.

2. Creating models

In this part, I will use 10 fold cross-validation to find a better model.

```
folds <- createFolds(y=census_train_new$income50,k=10)
```

(1) Logistic Regression / LDA / KNN

```
min=100
num=0
for(i in 1:10){

  fold_test <- census_train_new[folds[[i]],]
  fold_train <- census_train_new[-folds[[i]],]

  log.reg = glm((income50 == 1) ~
    # Quantitative
    age + fnlwgt + hours.per.week + education.num + capital.gain +
    capital.loss+

    # Qualitative
    sex + race + marital.status + native.country + workclass + occupation,
    data=fold_train,family=binomial)

  fold_predict <- predict(log.reg,type='response',newdata=fold_test)
  fold_predict = ifelse(fold_predict>0.5,1,0)
  fold_accuracy = mean(fold_predict != fold_test$income50)

  if(fold_accuracy<min)
  {
    min=fold_accuracy
    num=i
  }

}
```

```
## [1] 0.1392
```

Well, the original logistic regression model 's performance is not bad.

```
summary(log.reg)
```

```
##
## Call:
## glm(formula = (income50 == 1) ~ age + fnlwgt + hours.per.week +
##     education.num + capital.gain + capital.loss + sex + race +
##     marital.status + native.country + workclass + occupation,
##     family = binomial, data = fold_train)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -4.2844  -0.5017  -0.2037  -0.0467   3.7888
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    -1.191e+01  3.375e+01  -0.353  0.724075
## age             2.723e-02  1.922e-03  14.168 < 2e-16 ***
## fnlwgt          7.123e-07  2.061e-07   3.456  0.000548 ***
## hours.per.week  3.098e-02  1.940e-03  15.964 < 2e-16 ***
## education.num   2.836e-01  1.098e-02  25.829 < 2e-16 ***
## capital.gain    3.267e-04  1.247e-05  26.209 < 2e-16 ***
## capital.loss    6.677e-04  4.470e-05  14.936 < 2e-16 ***
## sex Male        1.277e-01  6.268e-02   2.037  0.041643 *
## race Asian-Pac-Islander 5.519e-01  2.910e-01   1.896  0.057921 .
## race Black      3.947e-01  2.743e-01   1.439  0.150197
## race Other     -8.768e-02  4.133e-01  -0.212  0.831999
## race White      5.318e-01  2.611e-01   2.037  0.041677 *
## marital.status.L -1.005e+00  1.336e-01  -7.522  5.39e-14 ***
## marital.status.Q -6.685e-02  1.616e-01  -0.414  0.679046
## marital.status.C  1.310e+00  1.434e-01   9.135 < 2e-16 ***
## marital.status^4 -1.355e+00  1.503e-01  -9.017 < 2e-16 ***
## marital.status^5  3.646e-01  1.842e-01   1.979  0.047761 *
## native.country.L  1.679e-01  6.461e-02   2.598  0.009364 **
## workclass.L     -4.632e+00  9.517e+01  -0.049  0.961182
## workclass.Q     -4.013e+00  9.517e+01  -0.042  0.966365
## workclass.C     -9.389e+00  1.163e+02  -0.081  0.935660
## workclass^4     -1.646e+00  4.970e+01  -0.033  0.973581
## workclass^5      1.545e+00  8.524e+01   0.018  0.985542
## workclass^6     -6.181e+00  1.257e+02  -0.049  0.960786
## workclass^7      2.971e+00  8.118e+01   0.037  0.970807
## occupation.L    -1.883e-02  8.730e-01  -0.022  0.982791
## occupation.Q     1.135e+00  8.894e-01   1.277  0.201700
## occupation.C    -5.120e-01  3.511e-01  -1.458  0.144799
## occupation^4    -1.926e+00  6.775e-01  -2.843  0.004468 **
## occupation^5    -3.123e-01  9.617e-01  -0.325  0.745354
## occupation^6     1.103e+00  1.133e+00   0.973  0.330316
## occupation^7     5.391e-01  1.184e+00   0.455  0.648904
## occupation^8    -2.351e-01  8.695e-01  -0.270  0.786900
## occupation^9    -1.106e+00  9.223e-01  -1.199  0.230461
## occupation^10     1.764e-01  3.518e-01   0.501  0.616052
## occupation^11     2.231e+00  8.143e-01   2.740  0.006152 **
## occupation^12    -7.552e-02  3.619e-01  -0.209  0.834718
## occupation^13    -1.045e+00  7.795e-01  -1.340  0.180190
## occupation^14    -1.499e+00  1.045e+00  -1.435  0.151382
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 24778  on 22499  degrees of freedom
## Residual deviance: 14407  on 22461  degrees of freedom
## AIC: 14485
##
## Number of Fisher Scoring iterations: 12
```


LDA

```
library(MASS)
min=100
num=0
for(i in 1:10){
  fold_test <- census_train_new[folds[[i]],]
  fold_train <- census_train_new[-folds[[i]],]
  lda1 = lda((income50 == 1) ~ age + workclass + fnlwgt + education.num +
             marital.status + occupation + race + sex + capital.gain +
             capital.loss + hours.per.week + native.country,

             data=fold_train)

  fold_predict <- predict(lda1,type='response',newdata=fold_test)
  fold_predict = ifelse(fold_predict$class == "TRUE", 1,0)
  fold_accuracy = mean(fold_predict != fold_test$income50)

  if(fold_accuracy<min)
  {
    min=fold_accuracy
    num=i
  }
}

print(min)
```

```
## [1] 0.1512
```

Little worse than logistic regression

KNN k=1,3,5,10,50,100

```
Knn_cv <- function(n,data) {

  min=100
  num=0
  for(i in 1:10){
    fold_test <- data[folds[[i]],]
    fold_train <- data[-folds[[i]],]

    train.X = cbind(fold_train$age, fold_train$fnlwgt,
                    fold_train$capital.gain, fold_train$capital.loss,
                    fold_train$hours.per.week, fold_train$native.country,
                    fold_train$workclass, fold_train$education.num,
                    fold_train$marital.status, fold_train$occupation,
                    fold_train$race, fold_train$sex)

    query.X = cbind(fold_test$age, fold_test$fnlwgt,
                    fold_test$capital.gain, fold_test$capital.loss,
                    fold_test$hours.per.week, fold_test$native.country,
                    fold_test$workclass, fold_test$education.num,
                    fold_test$marital.status, fold_test$occupation,
                    fold_test$race, fold_test$sex)

    set.seed(1)
    knn1 = knn(train.X, query.X, fold_train$income50, k = n)

    fold_accuracy = mean(knn1 != fold_test$income50)

    if(fold_accuracy<min)
    {
      min=fold_accuracy
      num=i
    }
  }

  print(min)
  print(num)

}
```

```
print("K = 1")
```

```
## [1] "K = 1"
```

```
Knn_cv(1,census_train_new)
```

```
## [1] 0.2524  
## [1] 9
```

```
print("K = 3")
```

```
## [1] "K = 3"
```

```
Knn_cv(3,census_train_new)
```

```
## [1] 0.2252  
## [1] 9
```

```
print("K = 5")
```

```
## [1] "K = 5"
```

```
Knn_cv(5,census_train_new)
```

```
## [1] 0.2084  
## [1] 9
```

```
print("K = 10")
```

```
## [1] "K = 10"
```

```
Knn_cv(10,census_train_new)
```

```
## [1] 0.2004  
## [1] 9
```

```
print("K = 50")
```

```
## [1] "K = 50"
```

```
Knn_cv(50,census_train_new)
```

```
## [1] 0.1956  
## [1] 9
```

```
print("K = 100")
```

```
## [1] "K = 100"
```

```
Knn_cv(100,census_train_new)
```

```
## [1] 0.2032  
## [1] 9
```

With K increase, the error rate is decrease, but if K keeps increasing, the error rate will increase. So I guess there always is a threshold value of K in KNN model.

The LDA and KNN both performance worse than logistic. We decided to go with the logistic regression and fine-tune it.

(2) Adjusting Logistic Regression model /improve models' performance

delete capital.loss and capital.gain

```
min=100
num=0
for(i in 1:10){
  fold_test <- census_train_new[folds[[i]],]
  fold_train <- census_train_new[-folds[[i]],]

  log.reg1 = glm((income50 == 1) ~
    # Quantitative
    age + fnlwgt + hours.per.week + education.num +

    # Qualitative
    sex + race + marital.status + native.country + workclass +
    occupation, data=fold_train, family=binomial)

  log.reg.probabilities1 = predict(log.reg1, fold_test, type="response")
  log.reg.predictions1 = ifelse(log.reg.probabilities1 > 0.5, 1, 0)
  fold_accuracy = mean(log.reg.predictions1 != fold_test$income50)

  if(fold_accuracy<min)
  {
    min=fold_accuracy
    num=i
  }
}

print(min)
```

```
## [1] 0.1524
```

The model performances worse before removing capital.loss and capital.gain, so I will keep use cubic splines with 5 nodes for all five quantitative predictors them.

```
min=100
num=0
for(i in 1:10){
  fold_test <- census_train_new[folds[[i]],]
  fold_train <- census_train_new[-folds[[i]],]

  log.reg2 = glm((income50 == 1) ~
    # Quantitative
    rcs(age, 5) + rcs(fnlwgt, 5) + rcs(capital.gain, 5) +
    rcs(capital.loss, 5) + rcs(hours.per.week, 5) + rcs(education.num, 5)
    +
    # Qualitative
    sex + race + marital.status + native.country + workclass +
    occupation,
    data=fold_train, family=binomial)

  log.reg.probabilities2 = predict(log.reg2, fold_test, type="response")
  log.reg.predictions2 = ifelse(log.reg.probabilities2 > 0.5, 1, 0)
  fold_accuracy = mean(log.reg.predictions2 != fold_test$income50)

  if(fold_accuracy<min)
  {
    min=fold_accuracy
    num=i
  }
}

print(min)
```

```
## [1] 0.156
```

reduce the number of nodes for fnlwgt and capital.loss from 5 to 3 and added several interaction terms

```
min=100
num=0
for(i in 1:10){
  fold_test <- census_train_new[folds[[i]],]
  fold_train <- census_train_new[-folds[[i]],]

  log.reg3 = glm((income50 == 1) ~

    # Quantitative
    rcs(age, 5) + rcs(fnlwgt, 3) + rcs(capital.gain, 5) +
    rcs(capital.loss, 3) + rcs(hours.per.week, 5) + rcs(education.num,5) +

    # Qualitative
    sex + race + marital.status + workclass + native.country
    + occupation +

    # Interactions
    marital.status * age + workclass * age +
    education.num * age + hours.per.week * workclass +
    hours.per.week * occupation + age * hours.per.week,
    data=fold_train, family=binomial)

  log.reg.proBABILITIES3 = predict(log.reg3, fold_test, type="response")
  log.reg.predictions3 = ifelse(log.reg.proBABILITIES3 > 0.5, 1, 0)
  fold_accuracy = mean(log.reg.predictions3 != fold_test$income50)

  if(fold_accuracy<min)
  {
    min=fold_accuracy
    num=i
  }
}
```

```
## [1] 0.1544
```

```
min=100
num=0
for(i in 1:10){
  fold_test <- census_train_new[folds[[i]],]
  fold_train <- census_train_new[-folds[[i]],]

  log.reg4 = glm((income50 == 1) ~

    # Quantitative
    rcs(age, 5) + fnlwgt + rcs(capital.gain, 5) +
    rcs(capital.loss, 5) + rcs(hours.per.week, 5) + rcs(education.num,5) +

    # Qualitative
    sex + race + marital.status + workclass + native.country
    + occupation +

    # Interactions
    marital.status * age + workclass * age +
    education.num * age + hours.per.week * workclass +
    hours.per.week * occupation + age * hours.per.week,
    data=fold_train, family=binomial)

  log.reg.proBABILITIES4 = predict(log.reg4, fold_test, type="response")
  log.reg.predictions4 = ifelse(log.reg.proBABILITIES4 > 0.5, 1, 0)
  fold_accuracy = mean(log.reg.predictions4 != fold_test$income50)

  if(fold_accuracy<min)
  {
    min=fold_accuracy
    num=i
  }
}
```

```
## [1] 0.1568
```

After trying these model, the performance still isn't satisfactory. I'd like to try tree.

(3)Tree / RandomForest

tree

```
min=100
num=0
for(i in 1:10){
  fold_test <- census_training[folds[[i]],]
  fold_train <- census_training[-folds[[i]],]

  library(tree)

  tree = tree(High_train~ age + fnlwgt + hours.per.week + education.num +
    capital.gain + capital.loss+
    # Qualitative
    sex + race + marital.status + workclass + occupation
    +native.country,
    data=fold_train)

  tree.pred = predict(tree, fold_test, type = "class")
  score = table(tree.pred, fold_test$High_train)
  fold_accuracy = (score[1,2] + score[2,1])/(score[1,2] + score[2,1] + score[1,1] +
    score[2,2])

  if(fold_accuracy<min)
  {
    min=fold_accuracy
    num=i
  }
}

print(min)
```

```
## [1] 0.1468
```

```
summary(tree)
```

```
##
## Classification tree:
## tree(formula = High_train ~ age + fnlwgt + hours.per.week + education.num +
##       capital.gain + capital.loss + sex + race + marital.status +
##       workclass + occupation + native.country, data = fold_train)
## Variables actually used in tree construction:
## [1] "marital.status" "capital.gain"  "education.num"
## Number of terminal nodes:  7
## Residual mean deviance:  0.6991 = 15730 / 22490
## Misclassification error rate: 0.156 = 3509 / 22500
```

Decision Tree performances better than LDA and KNN, but still worse than Logistic Regression

RandomForest

```
min=100
num=0
for(i in 1:10){
  fold_test <- census_training[folds[[i]],]
  fold_train <- census_training[-folds[[i]],]

  set.seed(3)
  library(randomForest)
  X = fold_train[,c(1:12)]
  y = fold_train[,15]
  rf = randomForest(X,y,mtry = 10,ntree = 100)

  rf.pred = predict(rf, fold_test, type = "class")
  score = table(rf.pred, fold_test$High_train)
  fold_accuracy = (score[1,2] + score[2,1])/(score[1,2] + score[2,1] + score[1,1] +
                                                score[2,2])

  if(fold_accuracy<min)
  {
    min=fold_accuracy
    num=i
  }
}
```

```
## [1] 0.1304
```

```
summary(rf)
```

```
##           Length Class  Mode
## call              5  -none- call
## type              1  -none- character
## predicted        22500 factor numeric
## err.rate          300  -none- numeric
## confusion          6  -none- numeric
## votes            45000 matrix numeric
## oob.times         22500 -none- numeric
## classes           2  -none- character
## importance         12  -none- numeric
## importanceSD        0  -none- NULL
## localImportance     0  -none- NULL
## proximity          0  -none- NULL
## ntree              1  -none- numeric
## mtry               1  -none- numeric
## forest            14  -none- list
## y                 22500 factor numeric
## test              0  -none- NULL
## inbag              0  -none- NULL
```

According to all models, I think the best one is randomForest, I will choose this one to do test.

3. Test Data

Do the same reprocessing as training data, and use randomForest to do the test.

```
set.seed(0)
test_rf.pred = predict(rf, census_test, type = "class")
score = table(test_rf.pred, census_test$High_train)
fold_accuracy = (score[1,2] + score[2,1]) / (score[1,2] + score[2,1] + score[1,1] +
                                              score[2,2])

fold_accuracy
```

```
## [1] 0.1596338
```

The final error rate is around 0.16. Not bad!

After doing this midterm project, I think preprocessing is very important. We have to view the data by eyes firstly. If we input the data directly, some missing data won't show off. That may cause big problem in fitting data and testing data.

And I think I should do more data processing, I should continue merging some variables. Although the best model in my project is randomForest, Logistic Regression also performances well. Logistic Regression is a basic classifier, but it is still very useful, and Logistic Regression can do many extensions.